

Implementing a basic driving agent.

When implementing the basic driving agent by initially defining action as:

```
action = random.choice((None, 'forward', 'left', 'right'))
```

The agent would sometimes find its destination and occasionally not (within an extended period of time before terminating the program) find its destination. Furthermore, using simple random choices, the agent would not find its destination within the deadline. The action definition was updated as suggested in a number of forum posts:

```
action = self.next_waypoint
```

With this definition, the agent always finds its destination. The agent seems to find the destination within the deadline and almost always with a positive reward at the end of each trial (this is based on observations of 30 trials.)

Definition of state

The state is defined as the inputs. The state consists of 'light', 'oncoming', 'right', and 'left'. These are the necessary inputs for the agent to make a decision on what action to take. Location was not included as the actual location is not relevant to what action to take since the actions are solely governed by what color the light is, whether there is oncoming traffic, and if it is clear to turn left or right. Finally, deadline is not included in state since the deadline is not *necessarily* required for the agent to take an action.

Although deadline can certainly be a component of state, I chose to leave it out because in my initial understanding of deadline was that it does not inform the action of the agent. However, deadline can in fact inform the action of the agent. The agent can be trained to find its destination within the deadline regardless of the rules of the road. Ultimately reason why I decided to not include deadline is because I am attempting to train an agent that adheres to the rules of the road rather than getting to the destination as quickly as possible. Although deadline can be relevant to the action of the agent, I chose to prioritize arriving to the destination with the closest adherence to the rules of the road as possible.

Implement Q-Learning

When q-learning was initially implemented, the agent took an extremely long time to reach its destination. The agent would repeatedly get stuck in suboptimal actions since it was looking for the max of similar states.

When state was updated to include next_waypoint the agent began to find reach its destination much more quickly. However, the agent did encounter an issue with red lights. The agent, after next_waypoint was included in state, repeatedly chose to go forward even if the light was red. The agent repeatedly attempts to go forward when in the following state, {'light': 'red', 'oncoming': None, 'right': None, 'left': None}. The agent is attempting to get to the destination so it is not 'confused', moving in circles, or moving away from the destination. The agent seems to frequently disregard the rule of going forward when the light is red. Simply put, the agent often acts to move towards the destination regardless of the rules when it is in the above state.

Enhancing the Driving Agent

In an effort to improve the driving agent I updated gamma in an effort to minimize its effect on the agents actions. Since over time the optimal action in each state will be known, I want to minimize gamma in order to have the agent use this known optimal action. In order to do so, gamma is initialized at 1 and over time the number of trials decreases gamma so its impact on the agents' actions are minimized. Gamma is decreased by the following definition:

```
gamma = 1 / len(self.q_table)
```

This was inspired by a number of discussions in the forums.¹

Based on the same reasoning as gamma above, I then updated the definition of alpha in order to minimize the impact of the learning rate, as the optimal actions become known over time. I gave alpha the same definition as gamma above:

```
alpha = 1 / len(self.q_table)
```

With the updated alpha & gamma it seems as though the agent experiences a bit of an improvement over time. The agent seems to take less time to reach the destination. However, it is not entirely clear whether or not the agent is finding the optimal action in certain states (primarily when the light is red and there is no oncoming traffic or other cars to the right or left.) That being said, the agent does seem to become more efficient over time.

To illustrate more concretely the agents actions over time I will use some examples of its actions from the last ten trials of the $\alpha=1/\text{len}(\text{self.q_table})$ & $\gamma=1$ trials (these can be found in the gamma_1.txt file.) In the last ten trials of this set of trial runs, the agent the rules of the road 32 times. In all of these instances the agent was in the {'light': 'red', 'oncoming': None,

¹ <https://discussions.udacity.com/t/question-on-the-maximization-step/170555/5>

'right': None, 'left': None}' state. In these instances the agent either chose to go forward (21/32 times) or left (11/32). There were 37 instances where the agent found itself in the above state; the agent chose the legal action in only 5 of those instances. The agent is attempting to reach its destination as quickly and as directly as possible regardless of the rules of the road. If the destination is ahead of the agent it will attempt to move forward even if the light is red.

Below is a table where I tuned the alpha & gamma parameters. An interesting note about this table is that high pass rate of the various values used for alpha & gamma.

Learning Rate	Discount Rate	Pass Rate	Mean Reward	Failed Trials
1	1/len(self.q_table)	100%	15	N/A
0.1	1/len(self.q_table)	100%	16.51	N/A
0.2	1/len(self.q_table)	100%	15.33	N/A
0.3	1/len(self.q_table)	100%	16.61	N/A
0.4	1/len(self.q_table)	100%	14.39	N/A
0.5	1/len(self.q_table)	100%	16.76	N/A
0.6	1/len(self.q_table)	99%	16.2626	23
0.7	1/len(self.q_table)	100%	15.96	N/A
0.8	1/len(self.q_table)	99%	15.58585	56
0.9	1/len(self.q_table)	100%	16.04	N/A
1/len(self.q_table)	1/len(self.q_table)	100%	14.32	N/A
1/len(self.q_table)	1	100%	16.55	N/A
1/len(self.q_table)	0.1	100%	15.34	N/A
1/len(self.q_table)	0.2	100%	15.1	N/A
1/len(self.q_table)	0.3	99%	16.1717	23
1/len(self.q_table)	0.4	100%	16.02	N/A
1/len(self.q_table)	0.5	100%	16.74	N/A
1/len(self.q_table)	0.6	99%	16.505	87
1/len(self.q_table)	0.7	99%	14.979	54
1/len(self.q_table)	0.8	100%	15.57	N/A
1/len(self.q_table)	0.9	100%	15.03	N/A

Additional Resources Used:

Reinforcement Learning 3 - Q Learning

<https://www.youtube.com/watch?v=1XRahNzA5bE>

Reinforcement Learning 4 - Q-Learning Parameters

<https://www.youtube.com/watch?v=XrxgdpduWOU>

Path Finding Q-Learning Tutorial

<http://mnemstudio.org/path-finding-q-learning-tutorial.htm>